

# 01\_exploration

January 20, 2026

## 1 01\_exploration.ipynb - Exploratory Data Analysis

**Objective:** Predict medical charges (`charges`) using regression models.

This notebook follows the following required workflow: 1. Load and inspect data 2. Check for missing values 3. Visualize distributions 4. Correlation analysis 5. Identify potential issues

### 2 1. Load and Inspect Data

#### 2.1 1.1 Import libraries

Import core python libraries for data analysis.

```
[ ]: # Load libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

#### 2.2 1.2 Load dataset

The dataset is loaded from `data/` folder to ensure reproducibility.

```
[ ]: # Load dataset from local path (repo data folder)
from pathlib import Path
DATA_PATH = Path("../data/insurance.csv")
df = pd.read_csv(DATA_PATH)

print("Insurance Dataset:")
print("Samples:", df.shape[0])
print("Features:", df.shape[1])
print("List of features:", df.columns.to_list())
df.head()
```

```
Insurance Dataset:
Samples: 1338
Features: 7
List of features: ['age', 'sex', 'bmi', 'children', 'smoker', 'region',
'charges']
```

```
[ ]:   age      sex      bmi  children smoker      region      charges
 0    19  female  27.900        0     yes southwest  16884.92400
 1    18     male  33.770        1     no southeast  1725.55230
 2    28     male  33.000        3     no southeast  4449.46200
 3    33     male  22.705        0     no northwest 21984.47061
 4    32     male  28.880        0     no northwest  3866.85520
```

### 2.3 1.3 Handle missing values

We check the dataset for any missing values.

```
[ ]: # Check for missing values
missing_values = df.isnull().sum()
print("Missing values:")
print(missing_values.to_string())
print("% of dataset with missing values:", missing_values.sum() / df.shape[0] * 100)
```

```
Missing values:
age          0
sex          0
bmi          0
children     0
smoker       0
region       0
charges      0
% of dataset with missing values: 0.0
```

## 3 2. Distribution Analysis

### 3.1 2.1. Age Distribution

We group samples by age groups to get raw values and percentages.

```
[ ]: # Visualize age distribution
plt.figure(figsize=(10, 6))
sns.histplot(df['age'], bins=20, kde=True)
plt.title('Age Distribution')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.show()

# Group samples by age groups - necessary for readability
age_groups = ['18-24', '25-30', '30-35', '35-40', '40-45', '45-50', '50-55', '55-60', '60-65']
df['age_group'] = pd.cut(df['age'], bins=[18, 25, 30, 35, 40, 45, 50, 55, 60, 65], labels=age_groups)
```


```

# Print size of age groups in both raw values and percentages
print("Age groups:")
print(df['age_group'].value_counts().to_string())

print("Age group percentages:")
print(df['age_group'].value_counts(normalize=True).to_string())

# Visualize age group distributions
plt.figure(figsize=(10, 6))
sns.countplot(data=df, x='age_group')
plt.title('Age Group Distribution')
plt.xlabel('Age Group')
plt.ylabel('Count')
plt.show()

```




Age groups:

age_group	Frequency
18-24	237
25-30	138
30-35	130
35-40	127
40-45	137
45-50	144
50-55	140
55-60	58
60-65	48
65-70	78
70-75	85
75-80	80
80-85	70
85-90	65
90-95	55
95-100	55
100-105	58
105-110	58
110-115	58
115-120	58
120-125	58
125-130	58
130-135	58
135-140	58
140-145	58
145-150	58
150-155	58
155-160	58
160-165	58
165-170	58
170-175	58
175-180	58
180-185	58
185-190	58
190-195	58
195-200	58
200-205	58
205-210	58
210-215	58
215-220	58
220-225	58
225-230	58
230-235	58
235-240	58
240-245	58
245-250	58
250-255	58
255-260	58
260-265	58
265-270	58
270-275	58
275-280	58
280-285	58
285-290	58
290-295	58
295-300	58
300-305	58
305-310	58
310-315	58
315-320	58
320-325	58
325-330	58
330-335	58
335-340	58
340-345	58
345-350	58
350-355	58
355-360	58
360-365	58
365-370	58
370-375	58
375-380	58
380-385	58
385-390	58
390-395	58
395-400	58
400-405	58
405-410	58
410-415	58
415-420	58
420-425	58
425-430	58
430-435	58
435-440	58
440-445	58
445-450	58
450-455	58
455-460	58
460-465	58
465-470	58
470-475	58
475-480	58
480-485	58
485-490	58
490-495	58
495-500	58
500-505	58
505-510	58
510-515	58
515-520	58
520-525	58
525-530	58
530-535	58
535-540	58
540-545	58
545-550	58
550-555	58
555-560	58
560-565	58
565-570	58
570-575	58
575-580	58
580-585	58
585-590	58
590-595	58
595-600	58
600-605	58
605-610	58
610-615	58
615-620	58
620-625	58
625-630	58
630-635	58
635-640	58
640-645	58
645-650	58
650-655	58
655-660	58
660-665	58
665-670	58
670-675	58
675-680	58
680-685	58
685-690	58
690-695	58
695-700	58
700-705	58
705-710	58
710-715	58
715-720	58
720-725	58
725-730	58
730-735	58
735-740	58
740-745	58
745-750	58
750-755	58
755-760	58
760-765	58
765-770	58
770-775	58
775-780	58
780-785	58
785-790	58
790-795	58
795-800	58
800-805	58
805-810	58
810-815	58
815-820	58
820-825	58
825-830	58
830-835	58
835-840	58
840-845	58
845-850	58
850-855	58
855-860	58
860-865	58
865-870	58
870-875	58
875-880	58
880-885	58
885-890	58
890-895	58
895-900	58
900-905	58
905-910	58
910-915	58
915-920	58
920-925	58
925-930	58
930-935	58
935-940	58
940-945	58
945-950	58
950-955	58
955-960	58
960-965	58
965-970	58
970-975	58
975-980	58
980-985	58
985-990	58
990-995	58
995-1000	58

```
55-60      125
60-65      91
Age group percentages:
```

```
age_group
18-24      0.186761
45-50      0.113475
50-55      0.110323
25-30      0.108747
40-45      0.107959
30-35      0.102443
35-40      0.100079
55-60      0.098503
60-65      0.071710
```



### 3.2 2.2 Sex Distribution

```
[ ]: # Print distribution of the sexes in raw values and percentages
print("Sex counts:")
print(df['sex'].value_counts().to_string())

print("Sex percentages:")
print(df['sex'].value_counts(normalize=True).to_string())

# Visualize sex distribution
```

```

plt.figure(figsize=(8, 6))
sns.countplot(data=df, x='sex')
plt.title('Sex Distribution')
plt.xlabel('Sex')
plt.ylabel('Count')
plt.show()

```

Sex counts:

```

sex
male      676
female    662


```

Sex percentages:

```

sex
male      0.505232
female    0.494768

```



#### 4 2.3. Body Mass Index (BMI) distribution

We group individual BMI into their respective groups to analyze how many of our patients would fall under one group.


```
[ ]: # Group samples into body mass index (bmi) groups - necessary for consistency
# with bmi scale
bmi_groups = ['Underweight', 'Normal', 'Overweight', 'Obese']
df['bmi_group'] = pd.cut(df['bmi'], bins=[0, 18.5, 24.9, 29.9, np.inf], labels=bmi_groups)

# Visualize BMI distribution
plt.figure(figsize=(10, 6))
sns.histplot(df['bmi'], bins=20, kde=True)
plt.title('Body Mass Index (BMI) Distribution')
plt.xlabel('BMI')
plt.ylabel('Frequency')
plt.show()

# Print size of BMI groups in raw values and percentages
print("BMI groups:")
print(df['bmi_group'].value_counts().to_string())

print("BMI percentages:")
print(df['bmi_group'].value_counts(normalize=True).to_string())

# Visualize BMI group distribution
plt.figure(figsize=(8, 6))
sns.countplot(data=df, x='bmi_group')
plt.title('BMI Group Distribution')
plt.xlabel('BMI Group')
plt.ylabel('Count')
plt.show()
```



BMI groups:


bmi\_group

Obese	716
Overweight	380
Normal	221
Underweight	21

BMI percentages:

bmi\_group

Obese	0.535127
Overweight	0.284006
Normal	0.165172
Underweight	0.015695



#### 4.1 2.4. Number of Children Distribution

```
[ ]: # Print number of children in terms of raw values and percentages
print("Number of children values:")
print(df['children'].value_counts().to_string())

print("Number of children percentages:")
print(df['children'].value_counts(normalize=True).to_string())


# Visualize children distribution
plt.figure(figsize=(8,6))
sns.countplot(data=df, x='children')
plt.title('Children Distribution')
plt.xlabel('Number of Children')
plt.ylabel('Count')
plt.show()
```

Number of children values:

children

0 574

```
1    324
2    240
3    157
4     25
5     18
Number of children percentages:
children
0    0.428999
1    0.242152
2    0.179372
3    0.117339
4    0.018685
5    0.013453
```



## 4.2 2.5. Smoker Status Distribution

```
[ ]: # Print distribution of smoking status in raw values and percentages
print("Smoker values:")
print(df['smoker'].value_counts().to_string())
```

```

print("Smoker percentages:")
print(df['smoker'].value_counts(normalize=True).to_string())

# Visualize smoker distribution
plt.figure(figsize=(8, 6))
sns.countplot(data=df, x='smoker')
plt.title('Smoker Distribution')
plt.xlabel('Smoker')
plt.ylabel('Count')
plt.show()


```

Smoker values:

smoker	Count
no	1064
yes	274

Smoker percentages:

smoker	Percentage
no	0.795217
yes	0.204783



### 4.3 2.6. Region Distribution

```
[ ]: # Print region distribution in raw values and percentages
print("Region values:")
print(df['region'].value_counts().to_string())

print("Region percentages:")
print(df['region'].value_counts(normalize=True).to_string())


# Visualize region distribution
plt.figure(figsize=(10, 6))
sns.countplot(data=df, x='region')
plt.title('Region Distribution')
plt.xlabel('Region')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.show()
```

Region values:

```
region
southeast    364
southwest    325
northwest    325
northeast    324
```

Region percentages:

```
region
southeast    0.272048
southwest    0.242900
northwest    0.242900
northeast    0.242152
```



#### 4.4 2.7 Insurance Charges Distribution

We group individual charges into insurance brackets to make analysis easier.

```
[ ]: # Group dataset into insurance brackets
insurance_brackets = [
    'Under $10,000', '$10,000 - $20,000', '$20,000 - $30,000',
    '$30,000 - $40,000', '$40,000 - $50,000', '$50,000 - $60,000',
    'Over $60,000'
]

df['insurance_bracket'] = pd.cut(
    df['charges'],
    bins=[0, 10000, 20000, 30000, 40000, 50000, 60000, float("inf")],
    labels=insurance_brackets
)

# Visualize insurance charge distribution
plt.figure(figsize=(10, 6))
sns.histplot(data=df, x="age", bins=20, kde=True)

plt.title('Insurance Charge Distribution')
```

```


plt.xlabel('Insurance Charge')
plt.ylabel('Frequency')
plt.show()

# Print insurance bracket size in terms of raw values and percentages
print("Insurance bracket values:")
print(df['insurance_bracket'].value_counts().to_string())

print("Insurance bracket percentages:")
print(df['insurance_bracket'].value_counts(normalize=True).to_string())

# Print insurance bracket distributions
plt.figure(figsize=(10, 6))
sns.countplot(data=df, x='insurance_bracket')
plt.title('Insurance Bracket Distribution')
plt.xlabel('Insurance Bracket')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.show()

```



Insurance bracket values:


insurance_bracket	
Under \$10,000	712
\$10,000 - \$20,000	353

\$20,000 - \$30,000	111
\$30,000 - \$40,000	83
\$40,000 - \$50,000	72
\$50,000 - \$60,000	4
Over \$60,000	3

Insurance bracket percentages:

insurance\_bracket

Under \$10,000	0.532138
\$10,000 - \$20,000	0.263827
\$20,000 - \$30,000	0.082960
\$30,000 - \$40,000	0.062033
\$40,000 - \$50,000	0.053812
\$50,000 - \$60,000	0.002990
Over \$60,000	0.002242



### 5 3. Correlation Analysis

We calculate the correlation between every feature and charges.

### 5.1 3.1. Smoking Correlation


```
[ ]: correlations = {}

# Print correlation between smoking and insurance charges
print("Correlation between smoking and insurance charges:")
print(df['smoker'].map({'yes': 1, 'no': 0}).corr(df['charges']))
correlations['smoker'] = (df['smoker'].map({'yes': 1, 'no': 0}).
    corr(df['charges']))


# Visualize correlation between smoking and insurance charges
plt.figure(figsize=(8, 6))
sns.scatterplot(data=df, x='smoker', y='charges')
plt.title('Smoking and Insurance Charge Correlation')
plt.xlabel('Smoker')
plt.ylabel('Insurance Charge')
plt.show()
```

Correlation between smoking and insurance charges:

0.7872514304984779




### 5.2 3.2. BMI Correlation

```
[ ]: # Print correlation between body mass index and insurance charges
print("Correlation between body mass index and insurance charges:")
print(df['bmi'].corr(df['charges']))
correlations['bmi'] = (df['bmi'].corr(df['charges']))  
  
# Visualize correlation between body mass index and insurance charges
plt.figure(figsize=(10, 6))
sns.scatterplot(data=df, x='bmi', y='charges')
plt.title('Body Mass Index and Insurance Charge Correlation')
plt.xlabel('Body Mass Index')
plt.ylabel('Insurance Charge')
plt.show()
```

Correlation between body mass index and insurance charges:

0.19834096883362895



### 5.3 3.3 Age Correlation

```
[ ]: # Print correlation between age and insurance charge
print("Correlation between age and insurance charge:")
print(df['age'].corr(df['charges']))
correlations['age'] = (df['age'].corr(df['charges']))
```


```

# Visualize correlation between age and insurance charge
plt.figure(figsize=(10, 6))
sns.scatterplot(data=df, x='age', y='charges')
plt.title('Age and Insurance Charge Correlation')
plt.xlabel('Age')
plt.ylabel('Insurance Charge')
plt.show()

```

Correlation between age and insurance charge:

0.29900819333064754




#### 5.4 3.4 Number of Children Correlation

```
[ ]: # Print correlation between number of children and insurance charges
print("Correlation between number of children and insurance charges:")
print(df['children'].corr(df['charges']))
correlations['children'] = (df['children'].corr(df['charges']))

# Correlation between number of children and insurance charges
plt.figure(figsize=(8, 6))
sns.scatterplot(data=df, x='children', y='charges')
plt.title('Number of Children and Insurance Charge Correlation')
plt.xlabel('Number of Children')
plt.ylabel('Insurance Charge')
plt.show()
```

Correlation between number of children and insurance charges:  
0.0679982268479048



## 5.5 3.5 Region Correlation


```
[ ]: # Print correlation between region and insurance charges
print("Correlation between region and insurance charges:")
print(df['region'].map({'southwest':0, 'southeast':1, 'northwest':2, □
    ↪'northeast':3}).corr(df['charges']))
correlations['region'] = (df['region'].map({'southwest':0, 'southeast':1, □
    ↪'northwest':2, 'northeast':3}).corr(df['charges']))
```

```
# Visualize correlation between region and insurance charges
plt.figure(figsize=(8, 6))
sns.scatterplot(data=df, x='region', y='charges')
plt.title('Region and Insurance Charge Correlation')
plt.xlabel('Region')
plt.ylabel('Insurance Charge')
plt.show()
```

Correlation between region and insurance charges:

0.006208234909444512



## 6 3.6. Sex Correlation


```
[ ]: # Print correlation between sex and insurance charges
print("Correlation between sex and insurance charges:")
print(df['sex'].map({'male':0, 'female':1}).corr(df['charges']))
correlations['sex'] = (df['sex'].map({'male':0, 'female':1}).
    ↪corr(df['charges'])))

# Visualize correlation between sex and insurance charges
df.groupby("sex")["charges"].mean()

plt.figure(figsize=(6,5))
sns.boxplot(data=df, x="sex", y="charges")
plt.title("Charges by Sex")
plt.show()
```

Correlation between sex and insurance charges:

-0.05729206220202548



### 6.1 3.7. Most Correlated Features

A number close to 1 implies a strong correlation, a number close to 0 implies a weak correlation, and a number close to -1 implies an inverse correlation.

```
[ ]: # Sort features by most correlated to insurance charges
correlations = {k: v for k, v in sorted(correlations.items(), key=lambda item: item[1], reverse=True)}
print("Features sorted by most correlated to insurance charges:")
print(list(correlations.keys()))
print("Correlation values:")
print(list(correlations.values()))
```

Features sorted by most correlated to insurance charges:

```
['smoker', 'age', 'bmi', 'children', 'region', 'sex']
```

Correlation values:

```
[np.float64(0.7872514304984779), np.float64(0.29900819333064754),
np.float64(0.19834096883362895), np.float64(0.0679982268479048),
np.float64(0.006208234909444512), np.float64(-0.05729206220202548)]
```

## 7 4. Written Analysis

- No missing values

### 7.1 Distributions

- Age is skewed towards 18-24 age group, but only by 7% more than the second largest age group (45-50).
- Smoking is split across the sample at around an 80% no to 20% yes split.
- Body Mass Index is skewed towards a range between 28 to 35. About 54% of the sample would be considered obese.
- The number of children is skewed towards 0-1. About half of the sample do not have children covered by health insurance.
- Insurance charges are mostly under \$10,000. The distribution is right-skewed, so applying a log-transform to charges may improve regression performance.
- Sex and Region are almost equally distributed between all features.

### 7.2 Correlations

- Smoking has the highest correlation with insurance charges (0.78).
- Region has the least correlation with insurance charges (0.006).
- Sex shows a very weak relationship with charges (correlation -0.06). This suggests sex alone is not a strong predictor compared with smoking, age, and BMI.

### 7.3 Issues with modelling

- Age, BMI and Charges have large ranges
- Sex, Region, and Smoker are categorical