

# 02\_modeling

January 20, 2026

## 1 02\_modeling.ipynb — Regression Modeling

**Objective:** Predict medical insurance charges (`charges`) using regression models.

This notebook follows the required workflow: 1. Preprocessing pipeline

2. Feature engineering
3. Model training (3+ models)
4. Model evaluation
5. Results visualization
6. Final analysis

### 1.1 1. Preprocessing Pipeline

#### 1.1.1 1.1 Import libraries

We import core Python libraries and scikit-learn utilities for preprocessing, modeling, evaluation, and visualization.

```
[1]: import numpy as np
      import pandas as pd

      import matplotlib.pyplot as plt
      import seaborn as sns

      from pathlib import Path

      from sklearn.model_selection import train_test_split
      from sklearn.compose import ColumnTransformer
      from sklearn.pipeline import Pipeline
      from sklearn.preprocessing import OneHotEncoder, StandardScaler, PolynomialFeatures
      from sklearn.impute import SimpleImputer

      from sklearn.linear_model import LinearRegression, Ridge, Lasso, ElasticNet
      from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

#### 1.1.2 1.2 Load dataset

The dataset is loaded from the local repository `data/` folder to ensure reproducibility.

```
[2]: # Load dataset from local path (repo data folder)
DATA_PATH = Path("../data/insurance.csv")
df = pd.read_csv(DATA_PATH)

df.head(), df.shape
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900		0	yes	southwest 16884.92400
1	18	male	33.770		1	no	southeast 1725.55230
2	28	male	33.000		3	no	southeast 4449.46200
3	33	male	22.705		0	no	northwest 21984.47061
4	32	male	28.880		0	no	northwest 3866.85520,

(1338, 7))

```
[3]: out_dir = Path("../figs")
out_dir.mkdir(exist_ok=True)
```

### 1.1.3 1.3 Define features/target and create train/test split (80/20)

We separate the feature matrix (**X**) and target (**y**), then split the dataset into training and testing sets.

A fixed random seed is used for reproducibility.

```
[4]: # Reproducibility: keep split consistent across runs

np.random.seed(42)

target = "charges"
X = df.drop(columns=[target])
y = df[target]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

X_train.shape, X_test.shape
```

```
[4]: ((1070, 6), (268, 6))
```

### 1.1.4 1.4 Preprocessing pipeline

We create a preprocessing pipeline to handle different feature types:

- **Numeric features:** median imputation + scaling
- **Categorical features:** mode imputation + one-hot encoding

This allows models to train correctly and ensures consistent preprocessing across all models.

```
[5]: # Build preprocessing pipeline for numeric + categorical features

numeric_features = X.select_dtypes(include=["int64", "float64"]).columns.
    ↪tolist()
categorical_features = X.select_dtypes(include=["object"]).columns.tolist()

numeric_transformer = Pipeline(steps=[
    ("imputer", SimpleImputer(strategy="median")),
    ("scaler", StandardScaler())
])

categorical_transformer = Pipeline(steps=[
    ("imputer", SimpleImputer(strategy="most_frequent")),
    ("onehot", OneHotEncoder(handle_unknown="ignore"))
])

preprocess = ColumnTransformer(
    transformers=[
        ("num", numeric_transformer, numeric_features),
        ("cat", categorical_transformer, categorical_features),
    ]
)

numeric_features, categorical_features
```

```
[5]: (['age', 'bmi', 'children'], ['sex', 'smoker', 'region'])
```

## 1.2 2. Feature Engineering

The target variable (**charges**) is typically right-skewed. To reduce the impact of extreme values, we try a **log transformation** of the target:

- Train on `log1p(charges)`
- Convert predictions back to dollar scale using `expm1()`

### 1.2.1 Utility: Evaluation function

To compare models consistently, we use a helper function that trains a model and reports:

- **MAE** (Mean Absolute Error)
- **RMSSE** (Root Mean Squared Error)
- **R<sup>2</sup>** (Coefficient of Determination)

```
[6]: def evaluate_model(name, model, X_train, y_train, X_test, y_test):
    model.fit(X_train, y_train)
    preds = model.predict(X_test)

    mae = mean_absolute_error(y_test, preds)
    rmse = np.sqrt(mean_squared_error(y_test, preds))
```

```
r2 = r2_score(y_test, preds)

return {
    "Model": name,
    "MAE": mae,
    "RMSE": rmse,
    "R2": r2,
    "y_pred": preds
}
```

### 1.3 3. Model Training (3+ Models)

### 1.3.1 3.1 Linear Regression (baseline)

Linear Regression is used as a baseline model to understand the performance without regularization.

### 1.3.2 3.1 Linear Regression (baseline)

Linear Regression is used as a baseline model to understand performance without regularization.

```
[7]: # Baseline model: Linear Regression
```

```
model_lr = Pipeline(steps=[  
    ("preprocess", preprocess),  
    ("model", LinearRegression())  
])  
  
res_lr = evaluate_model("Linear Regression", model_lr, X_train, y_train,  
    X_test, y_test)  
res_lr
```

```
[7]: {'Model': 'Linear Regression',
      'MAE': 4181.194473753651,
      'RMSE': np.float64(5796.284659276274),
      'R2': 0.7835929767120722,
      'y_pred': array([ 8969.55027444,  7068.74744287, 36858.41091155,
9454.67850053,
26973.17345656, 10864.11316424, 170.28084136, 16903.45028662,
1092.43093614, 11218.34318352, 28101.68455267, 9377.73460205,
5263.0595179 , 38416.04221107, 40255.82339284, 37098.25353123,
15240.39392306, 35912.88264434, 9112.52398703, 31461.92108909,
3847.68845883, 10130.12001517, 2370.54189389, 7140.21550828,
11301.76782638, 12961.65366224, 14509.47251876, 6159.8976107 ,
9963.85857263, 2177.85718217, 9115.93673493, 13073.68932159,
4561.82376202, 3408.20756033, 4459.81359745, 13032.06505076,
1979.99357292, 8813.28303302, 33271.29124448, 32585.51583927,
3908.76090964, 4326.10774721, 14142.81326533, 11423.45494846,
8774.13955311, 12097.28051001, 5281.57353499, 3150.5596042 ,
```

35494.46461214, 9150.1124786 , 15836.84575621, 2343.57470069,  
 12364.78414194, 1482.29488266, 13389.06105161, 12573.57395972,  
 4341.83680558, 32165.33688042, 13321.3360032 , 12896.82071102,  
 14167.99421483, 10506.17623512, 16360.78543548, 7763.89824584,  
 11839.25019431, 4061.19750503, 26652.40230125, 10930.14138671,  
 2137.41385988, 6209.01123411, 10729.82391284, 11628.3104129 ,  
 10981.04528946, 9166.50818596, 11954.27732874, 6747.85121734,  
 7248.5304713 , 10735.16710748, 6580.84819774, 8762.00329355,  
 3767.13383454, 36632.4975496 , 6378.11979721, 30842.09248656,  
 34846.52451051, 35278.07387112, 7019.444352 , 12861.38414264,  
 9942.30149778, 14473.5260648 , 17693.37304474, 35258.24845137,  
 33029.58968269, 6185.91730447, 31999.98962535, 9481.33158273,  
 29444.04271523, 3674.48498404, 28308.26432106, 5823.36495229,  
 5407.76752001, 1883.4947576 , 11499.675042 , 15075.90690632,  
 11699.63163008, 4308.82427855, 9895.1840044 , 31708.40056201,  
 -86.87094667, 32819.71429004, 3280.69178415, 10183.88853878,  
 14318.76389179, 31642.35684542, 11461.57806791, 3929.23701831,  
 13107.89313088, 31810.99450607, 8152.02593593, 3238.08417076,  
 8439.56108376, 10594.63871458, 15219.68736374, 5647.8808143 ,  
 3781.95285499, 10228.944897 , 10900.12933883, 11122.74845192,  
 14438.14112575, 7430.31504776, 5386.22676759, 9231.32739901,  
 9343.76283713, 12538.27606344, 8337.66982683, 15333.36900871,  
 8411.2145439 , 31797.27496298, 35785.91843418, 31603.71967017,  
 6011.96229251, 12607.03584641, 6013.5115031 , 14560.79590559,  
 2493.47989441, 32963.45524228, 6265.14380504, 5034.62173797,  
 14344.81347407, 6941.1412259 , 38670.01270366, 3087.58741836,  
 5885.8752536 , 31686.24200595, 11562.61859836, 8476.04749512,  
 14806.72486264, 9814.46186143, 27105.71831469, 33453.83352069,  
 14551.8999207 , 1684.36856768, 13166.96197398, 2222.76894041,  
 5449.59393727, 11568.96325488, 39807.96912709, 36500.65163031,  
 34001.37945748, 3897.27856532, 7456.14132125, 8661.82084477,  
 12450.92458882, 4813.53293089, 2047.65528159, 32112.11251984,  
 25111.52085938, 17484.27663755, 26411.46181822, 10159.52421 ,  
 37260.32666386, -441.23918333, 6779.55013103, 7781.45337795,  
 4367.95988484, 5105.87170813, 5919.18675042, 4305.71645941,  
 15191.08806502, 11132.09935114, 6932.80116584, 2525.64793222,  
 1536.05183213, 31944.78284317, 16414.12251517, 12011.53367195,  
 1268.05926603, 12531.25953189, 1564.93415917, 8737.33621694,  
 1873.03940488, 33916.22971211, 10858.38635063, 2603.43633853,  
 25674.40250332, 26343.43022704, 9430.91152033, 1800.73500777,  
 13261.42480211, 1120.17810533, 10386.66427709, 10567.29006474,  
 16944.25995713, 26846.54662457, 6939.11178393, 5193.04710054,  
 5846.00017265, 13229.60536846, 11098.33930228, 8362.28134289,  
 5135.53940151, 12308.34064139, 13861.17886997, 35773.70926219,  
 4157.01930317, 28917.86562624, -914.37342357, 2873.71150671,  
 11046.2540774 , 15683.06950225, 5210.67532324, 6888.38518351,  
 3854.31140958, 31312.64705453, 7241.43226665, 12405.99508651,

```

5619.17039188, 9528.22557021, 36314.009043 , 4429.40596906,
9667.91523953, 31161.15738995, 5747.13292318, 4603.37294255,
1048.35533791, 4832.6604097 , 4574.9041044 , 6507.30666036,
18659.12407756, -1545.57184934, 2376.4352498 , 10694.62157146,
3151.28919904, 10209.96361187, 3733.89128353, 5125.08103172,
12400.90700504, 6218.65296628, 8231.63765089, 7590.50155269,
8924.15352268, 10482.90359975, 27808.04576398, 39061.50093248,
11761.4991981 , 7687.56363151, 40920.29151165, 12318.58665305])}

```

### 1.3.3 3.2 Ridge Regression (L2 regularization)

Ridge Regression adds L2 regularization to reduce overfitting risk and stabilize coefficients.

```
[8]: # Regularized model: Ridge Regression (L2)

model_ridge = Pipeline(steps=[
    ("preprocess", preprocess),
    ("model", Ridge(alpha=1.0, random_state=42))
])

res_ridge = evaluate_model("Ridge Regression", model_ridge, X_train, y_train, ▾
                           X_test, y_test)
res_ridge
```

```
[8]: {'Model': 'Ridge Regression',
      'MAE': 4186.913071783853,
      'RMSE': np.float64(5798.298795415483),
      'R2': 0.7834425531348179,
      'y_pred': array([ 8979.95466915,  7079.65062176, 36792.69449533,
9468.40383721,
26924.20409129, 10877.48262742, 189.33407024, 16906.5364521 ,
1110.77813969, 11230.05473287, 28054.55753557, 9391.18130892,
5282.67084845, 38361.2969438 , 40196.55588508, 37044.2255803 ,
15249.07493253, 35855.48707115, 9127.19755076, 31402.9798776 ,
3871.39372135, 10144.41490692, 2393.01899883, 7152.50240329,
11309.34227107, 12966.14684124, 14513.03575498, 6174.08278797,
9973.11162865, 2204.66791018, 9128.69677815, 13077.1762597 ,
4581.96986124, 3425.13611258, 4480.34545672, 13040.94839174,
1999.40018854, 8824.34880891, 33210.59589532, 32536.19394796,
3928.71585688, 4345.24339325, 14147.58951355, 11432.57403284,
8794.46068391, 12108.05348056, 5296.37626983, 3169.76370446,
35442.953054 , 9170.14128517, 15842.93358915, 2370.82998812,
12370.51925492, 1504.74757968, 13402.08831489, 12583.91919074,
4360.65412822, 32112.13205652, 13326.25478627, 12901.43408637,
14171.76154193, 10520.8309281 , 16368.20110073, 7774.99600127,
11845.02966696, 4076.30357816, 26607.88835336, 10947.82427071,
2154.70241403, 6227.28863654, 10742.11517691, 11635.86177334,
```

10993.0273339 , 9186.10458716, 11970.05806062, 6764.91585053,  
 7266.29856901, 10741.16969246, 6599.74221195, 8774.25761539,  
 3792.1663936 , 36571.85277991, 6392.94110238, 30786.81296008,  
 34795.5470087 , 35219.47875505, 7042.38897179, 12869.19585924,  
 9951.05230003, 14483.34518812, 17693.7042813 , 35200.04126331,  
 32970.14096714, 6203.97820341, 31951.44029089, 9496.41378992,  
 29398.23809593, 3695.61783252, 28255.17769525, 5837.60293229,  
 5421.35112831, 1902.48767119, 11515.47124405, 15084.93571556,  
 11709.08849608, 4334.00907279, 9913.75325025, 31658.5014513 ,  
 -62.76505139, 32765.26618271, 3302.20999859, 10202.468514 ,  
 14322.23886987, 31585.04424111, 11468.53026374, 3950.63921647,  
 13115.9182112 , 31754.66847562, 8168.79760916, 3258.31421508,  
 8450.4755043 , 10606.80334935, 15223.14367205, 5670.72434225,  
 3797.97227844, 10237.49748834, 10913.07524121, 11134.05494344,  
 14445.91013296, 7451.3325795 , 5404.93350914, 9251.68222709,  
 9359.38273613, 12543.68231595, 8348.82510042, 15340.73259083,  
 8425.58888995, 31743.77429624, 35722.36977395, 31545.67294629,  
 6023.97884981, 12612.05830817, 6035.8660371 , 14564.2138594 ,  
 2519.59485721, 32909.60100643, 6282.2849207 , 5049.67207707,  
 14349.1300196 , 6958.59185369, 38604.8159307 , 3103.35954641,  
 5908.7793907 , 31632.98427022, 11569.3026176 , 8490.30164481,  
 14812.856352 , 9828.34400546, 27058.43535884, 33398.29840324,  
 14555.74842927, 1702.44577746, 13176.9676734 , 2244.18205203,  
 5470.47698854, 11585.87286969, 39740.94173367, 36435.67537238,  
 33945.05990595, 3913.83850364, 7468.32556817, 8673.25375981,  
 12460.7315404 , 4831.68802593, 2066.31956912, 32062.71857768,  
 25066.25748762, 17492.57424545, 26359.65063862, 10166.08137342,  
 37194.97123094, -420.29378645, 6794.80366089, 7794.07326342,  
 4388.65903137, 5119.46557669, 5937.01329536, 4321.58662978,  
 15199.30565221, 11141.73295969, 6952.12158074, 2543.00408127,  
 1554.2430066 , 31892.93587814, 16423.6513693 , 12022.27436864,  
 1286.57507942, 12537.44445658, 1583.30247041, 8747.4996151 ,  
 1889.9354484 , 33858.23305605, 10876.66729281, 2625.16173483,  
 25629.19477356, 26296.53524992, 9442.59687648, 1819.77224804,  
 13275.00166848, 1140.28301568, 10398.46464543, 10579.86015035,  
 16946.21068677, 26803.17080046, 6961.00297225, 5206.97030764,  
 5864.78977892, 13240.50508575, 11109.5750739 , 8378.31613247,  
 5149.25768797, 12317.30582424, 13876.55212454, 35721.28435372,  
 4175.71394124, 28872.20377634, -892.43280956, 2893.91385101,  
 11062.62821909, 15692.52243283, 5234.1351444 , 6898.70001166,  
 3871.94735321, 31259.60507717, 7251.32324978, 12420.42130136,  
 5637.35531449, 9541.9681777 , 36257.06947093, 4453.63700689,  
 9675.56719544, 31110.23979821, 5759.99929807, 4623.83496176,  
 1070.48079037, 4850.43565873, 4591.83275741, 6525.37002991,  
 18665.90505616, -1520.71355895, 2394.85693335, 10709.09655715,  
 3168.47940181, 10217.20016585, 3749.80921397, 5145.04529563,  
 12410.68904626, 6235.58715117, 8245.37092732, 7603.70447919,

```
8934.7518229 , 10488.91250026, 27763.60505879, 39003.33603151,  
11766.31670566, 7703.06007585, 40856.63797972, 12325.54427666])}
```

### 1.3.4 3.3 Lasso Regression (L1 regularization)

Lasso Regression uses L1 regularization, which can shrink some coefficients to zero and act like feature selection.

```
[9]: # Regularized model: Lasso Regression (L1)
```

```
model_lasso = Pipeline(steps=[  
    ("preprocess", preprocess),  
    ("model", Lasso(alpha=0.001, random_state=42, max_iter=10000))  
])  
  
res_lasso = evaluate_model("Lasso Regression", model_lasso, X_train, y_train, □  
                           ↵X_test, y_test)  
res_lasso
```

```
[9]: {'Model': 'Lasso Regression',  
      'MAE': 4181.194836701594,  
      'RMSE': np.float64(5796.284902970282),  
      'R2': 0.7835929585152103,  
      'y_pred': array([ 8969.54880594, 7068.74574693, 36858.40184695,  
9454.67883166,  
                    26973.17048398, 10864.11961252, 170.28080622, 16903.45077161,  
1092.43275295, 11218.3430439 , 28101.68531136, 9377.73543993,  
5263.05729335, 38416.03426038, 40255.81728306, 37098.2466965 ,  
15240.38848613, 35912.87242864, 9112.52109327, 31461.91305972,  
3847.694268 , 10130.12635752, 2370.55124339, 7140.2151526 ,  
11301.7632743 , 12961.65251017, 14509.47490621, 6159.8939425 ,  
9963.85684931, 2177.86065665, 9115.93757845, 13073.68639255,  
4561.82214443, 3408.20488984, 4459.8218543 , 13032.0679869 ,  
1979.99804526, 8813.28603715, 33271.28228203, 32585.51074824,  
3908.76207369, 4326.11022984, 14142.8149818 , 11423.45704425,  
8774.14009029, 12097.28741414, 5281.57059989, 3150.56252672,  
35494.45750877, 9150.11282495, 15836.84191141, 2343.57746816,  
12364.78020763, 1482.30064951, 13389.06218253, 12573.57591011,  
4341.84205379, 32165.33067122, 13321.3387085 , 12896.81954576,  
14167.99681616, 10506.17168659, 16360.78971537, 7763.89522325,  
11839.24933871, 4061.19585987, 26652.40256966, 10930.14122085,  
2137.41273753, 6209.01943811, 10729.82670466, 11628.30811401,  
10981.04250969, 9166.50880119, 11954.27786184, 6747.84940296,  
7248.53271202, 10735.16490468, 6580.84485723, 8762.00507667,  
3767.13686073, 36632.49475289, 6378.12256313, 30842.08451037,  
34846.5182043 , 35278.07110358, 7019.44470389, 12861.38563369,  
9942.29780101, 14473.53068425, 17693.36932182, 35258.23968437,
```

33029.58247083, 6185.92068099, 31999.98462191, 9481.33806753,  
 29444.03850875, 3674.48362075, 28308.25491236, 5823.36416422,  
 5407.76550547, 1883.4998459 , 11499.67627324, 15075.9114327 ,  
 11699.63554797, 4308.8264069 , 9895.18455403, 31708.39699204,  
 -86.86466949, 32819.70528593, 3280.70083765, 10183.88903067,  
 14318.76178235, 31642.34963913, 11461.57657919, 3929.24524205,  
 13107.89655651, 31810.99073754, 8152.02239294, 3238.08704151,  
 8439.5600637 , 10594.63915671, 15219.68410253, 5647.88336594,  
 3781.95064752, 10228.94103846, 10900.12576088, 11122.75149299,  
 14438.13759872, 7430.31698545, 5386.22525893, 9231.32720294,  
 9343.75883801, 12538.27474057, 8337.67350164, 15333.37486043,  
 8411.2185715 , 31797.27689048, 35785.90903286, 31603.71323239,  
 6011.96094025, 12607.03470683, 6013.51400439, 14560.7910153 ,  
 2493.48350618, 32963.4557242 , 6265.14538098, 5034.62621278,  
 14344.80796662, 6941.14898427, 38670.00071351, 3087.58657157,  
 5885.87752219, 31686.23913644, 11562.61485638, 8476.05153591,  
 14806.72780433, 9814.45877367, 27105.71065659, 33453.83035053,  
 14551.89708232, 1684.36738702, 13166.96804851, 2222.77463784,  
 5449.60046316, 11568.96319163, 39807.96474957, 36500.64244213,  
 34001.37624636, 3897.28320436, 7456.1449724 , 8661.82376105,  
 12450.92750831, 4813.53556348, 2047.66041932, 32112.1079598 ,  
 25111.52417284, 17484.27671085, 26411.45374806, 10159.52222868,  
 37260.31927764, -441.2377608 , 6779.55267266, 7781.45640916,  
 4367.96313091, 5105.87023923, 5919.19079825, 4305.72113699,  
 15191.08322869, 11132.10140459, 6932.80720937, 2525.64894016,  
 1536.05329901, 31944.78303941, 16414.12286795, 12011.53065989,  
 1268.05822075, 12531.25721355, 1564.93547955, 8737.34013882,  
 1873.0390559 , 33916.22561264, 10858.38610715, 2603.43862545,  
 25674.39469187, 26343.42309107, 9430.91395565, 1800.73258585,  
 13261.42570033, 1120.1829944 , 10386.67249349, 10567.2870749 ,  
 16944.25592941, 26846.5437103 , 6939.11315782, 5193.04765231,  
 5845.99789563, 13229.60027511, 11098.33689569, 8362.27831608,  
 5135.53782014, 12308.34193527, 13861.17723675, 35773.70272112,  
 4157.0224939 , 28917.8621069 , -914.36742537, 2873.71501197,  
 11046.25514574, 15683.06261512, 5210.67791423, 6888.38406037,  
 3854.31489346, 31312.6491739 , 7241.43106296, 12405.99629721,  
 5619.17193797, 9528.23338027, 36313.99787234, 4429.40884842,  
 9667.91536428, 31161.15777249, 5747.13128001, 4603.37082971,  
 1048.3594368 , 4832.66351782, 4574.90731966, 6507.31445699,  
 18659.11567227, -1545.56690686, 2376.44021206, 10694.62697428,  
 3151.29418351, 10209.96345752, 3733.88943633, 5125.07871494,  
 12400.91449783, 6218.65497053, 8231.64228531, 7590.4993997 ,  
 8924.14940108, 10482.90165037, 27808.04257129, 39061.49531393,  
 11761.49673008, 7687.56741272, 40920.2805605 , 12318.58879141])}

### 1.3.5 3.4 Ridge Regression (log-transformed target)

This model applies Ridge Regression but trains on the log-transformed target to reduce skew effects.

```
[10]: # Ridge trained on log1p(target), then predictions converted back using expm1()

y_train_log = np.log1p(y_train)
y_test_log = np.log1p(y_test)

model_ridge_log = Pipeline(steps=[
    ("preprocess", preprocess),
    ("model", Ridge(alpha=1.0, random_state=42))
])

model_ridge_log.fit(X_train, y_train_log)
pred_log = model_ridge_log.predict(X_test)

pred_back = np.expm1(pred_log) # back to dollars scale

res_ridge_log = {
    "Model": "Ridge (log target)",
    "MAE": mean_absolute_error(y_test, pred_back),
    "RMSE": np.sqrt(mean_squared_error(y_test, pred_back)),
    "R2": r2_score(y_test, pred_back),
    "y_pred": pred_back
}

res_ridge_log
```

```
[10]: {'Model': 'Ridge (log target)',
'MAE': 3881.879523350427,
'RMSE': np.float64(7780.62105852155),
'R2': 0.6100575929052479,
'y_pred': array([ 9089.39360191,  5606.81781822,  65769.15161747,
9126.84667347,
        14034.7786963 ,  5948.65656835,  2830.505004 ,  15103.08553442,
3794.96124214, 10525.90538346, 22803.91885629,  7490.39093995,
4446.11944791, 49639.46792687, 59562.50812646, 44713.83886397,
11546.36167356, 42539.81003209,  7788.28858304, 32070.00515699,
4887.94944385, 7717.46043724,  2739.52762982,  4196.79840541,
11637.41336671, 11382.90637478, 12782.2667212 ,  5465.64905952,
9923.19408666, 2603.44296409,  8625.78953271, 11784.77331591,
3312.44350847, 5254.75160502,  3797.11378237,  8560.34027115,
3281.22050101, 7303.39263268, 46094.68102827, 26599.0769917 ,
4569.09144507, 3602.41672615, 12343.6834439 , 10749.13895121,
5646.40526353, 11348.43666109, 4240.65245299, 4398.8199929 ,
41079.90455084, 5465.58777929, 14207.48779643, 2825.79211556,
8232.37432172, 2822.62199599, 10125.09520552, 11005.68351029,
4217.14504486, 28866.5048732 , 11617.15841724, 11085.65822176,
13542.60295716, 6064.52535627, 16429.68841247, 7733.87746832,
10372.78129143, 4424.80300265, 18711.86709635, 10446.76662782,
```

3811.91929278, 3435.17811261, 7118.32681903, 10361.56080016,  
 8169.15184329, 6618.90689171, 7738.97033678, 5173.59364584,  
 5531.95736428, 10472.16260625, 4935.99845352, 8886.03306416,  
 2911.32181562, 60192.31263095, 5491.00379324, 23520.44553075,  
 20866.26167083, 42811.72524435, 5280.44181294, 10354.0818451 ,  
 8647.31684571, 12581.88058026, 19376.73440828, 57639.00737998,  
 42228.13205727, 5407.87490053, 24185.21692406, 7168.10137938,  
 21304.06052789, 3120.72797107, 25960.59762429, 6308.04508167,  
 4899.82532033, 2890.26503326, 6561.00495302, 13847.0933026 ,  
 11881.87209106, 3058.74867633, 7679.40475692, 35718.69528424,  
 2649.61188576, 43331.30140403, 2841.90152031, 4690.89257401,  
 14259.80472606, 29288.13496071, 10254.06808869, 3221.1548204 ,  
 12279.32701271, 36029.35957862, 6517.73495856, 3531.07687256,  
 6348.09819802, 7866.8902577 , 13817.23872345, 3455.02413363,  
 4717.28595037, 8340.33756071, 7792.14338303, 9401.25399375,  
 12207.04722771, 3624.87895733, 4451.70062419, 6503.90360527,  
 6519.88546281, 9028.7867699 , 6040.67830796, 12929.21903012,  
 5337.31123425, 28113.95889812, 39095.5684499 , 29134.31167306,  
 5373.03075015, 10446.61750223, 3677.64148488, 12910.0033073 ,  
 3115.95693863, 40474.86078835, 5459.63899482, 4369.24312008,  
 13473.20596152, 5201.30648461, 68209.60637482, 3502.43882331,  
 3023.16302177, 25999.71736568, 7599.24862174, 5480.43850694,  
 13791.77523229, 8205.11864366, 14190.59267035, 33787.33776177,  
 15519.39885412, 3243.83668848, 14392.05868707, 3276.51925979,  
 4276.38974283, 7832.63083875, 68924.41194536, 54861.48830932,  
 37998.09036493, 3616.97540435, 7978.40651186, 6598.47222166,  
 7247.28602869, 4553.58848385, 3125.28776329, 33923.63610221,  
 12408.28869247, 14451.96476013, 23396.55893686, 10478.97314119,  
 68023.97426305, 3254.88412679, 7219.24145773, 5652.56478655,  
 5166.24997097, 3890.98008378, 4731.03589378, 4142.75119017,  
 10190.28145559, 9657.08531299, 4539.82706902, 3167.41668446,  
 3348.25848615, 25289.64470171, 13197.27217216, 9147.12140909,  
 3189.85763189, 13004.2090908 , 3047.10408094, 7980.207993 ,  
 3582.50631404, 40783.07850724, 5717.14356062, 3373.25928711,  
 14790.70169091, 16653.81141574, 8003.60382604, 4069.57178806,  
 8323.53034666, 3486.74954581, 10553.25981163, 10735.75588997,  
 12836.96696327, 15423.5013897 , 6196.6707973 , 4695.72080676,  
 5007.46422494, 14571.48783937, 10927.3782748 , 5812.68738786,  
 3626.61808076, 7789.3757194 , 8220.96195182, 38533.33060757,  
 3318.5290882 , 17585.17504293, 2582.00464431, 2933.60967522,  
 8816.41806942, 13663.30412399, 3160.23972286, 7960.20370462,  
 4617.44657422, 28235.60894227, 8269.14534356, 8235.55859449,  
 4501.42383284, 6327.24844691, 41372.66932531, 3136.19337278,  
 11501.37010469, 23362.30565832, 4089.23974688, 4030.49418466,  
 2725.50687197, 3751.53298759, 4641.48194726, 4758.21790128,  
 17839.85642722, 2454.91226992, 3019.44161642, 7924.3817494 ,  
 3770.19531931, 10943.72748283, 3523.19157549, 4019.12775708,

```
10912.52099328, 4277.17319661, 7760.7303938 , 6638.61937309,
8330.52918915, 11679.02108235, 16419.76764442, 60973.44439393,
11447.55893701, 6296.9520442 , 57048.13209384, 10179.65177468])}
```

## 1.4 4. Model Evaluation

We compare models using **MAE**, **RMSE**, and **R<sup>2</sup>** on the test set.

### 1.4.1 4.1 Model comparison table

The table below summarizes performance for all trained models.

```
[11]: # Build model comparison table (lower RMSE/MAE is better)
```

```
results = pd.DataFrame([
    {k: res_lr[k] for k in ["Model", "MAE", "RMSE", "R2"]},
    {k: res_ridge[k] for k in ["Model", "MAE", "RMSE", "R2"]},
    {k: res_lasso[k] for k in ["Model", "MAE", "RMSE", "R2"]},
    {k: res_ridge_log[k] for k in ["Model", "MAE", "RMSE", "R2"]},
]).sort_values("RMSE")

results_round = results.copy()
results_round["MAE"] = results_round["MAE"].round(2)
results_round["RMSE"] = results_round["RMSE"].round(2)
results_round["R2"] = results_round["R2"].round(4)
results_round
```

```
[11]:
```

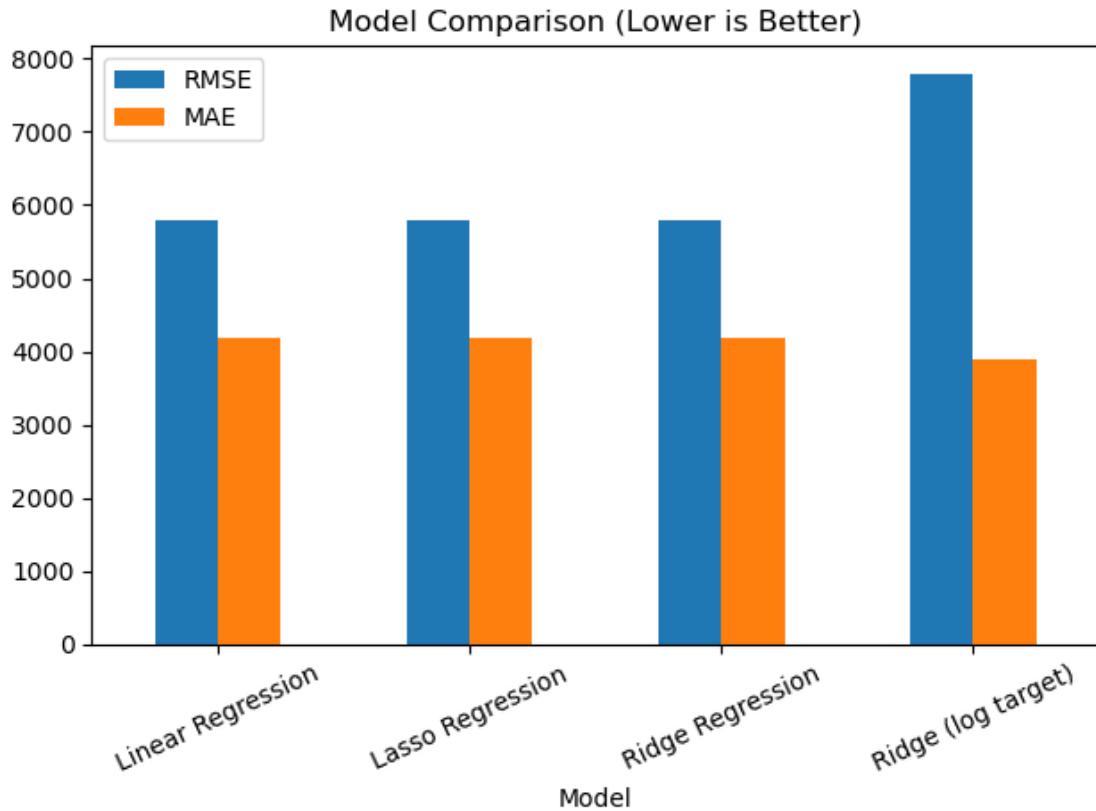
	Model	MAE	RMSE	R2
0	Linear Regression	4181.19	5796.28	0.7836
2	Lasso Regression	4181.19	5796.28	0.7836
1	Ridge Regression	4186.91	5798.30	0.7834
3	Ridge (log target)	3881.88	7780.62	0.6101

## 1.5 5. Results Visualization

We visualize the comparison results and inspect the best model using diagnostic plots.

### 1.5.1 5.1 Model Comparison (MAE & RMSE)

```
[12]: results.set_index("Model")[["RMSE", "MAE"]].plot(kind="bar", rot=25)
plt.title("Model Comparison (Lower is Better)")
plt.tight_layout()
plt.show()
```



### 1.5.2 5.2 Select best model

The best model is selected based on the lowest RMSE on the test set.

```
[13]: # Select best model based on lowest RMSE
best_model = results.iloc[0]["Model"]
print("Best model by RMSE:", best_model)
```

Best model by RMSE: Linear Regression

### 1.5.3 5.3 Predicted vs Actual plot (best model)

This scatter plot compares predicted charges to actual charges.  
A strong model should place points close to the diagonal line.

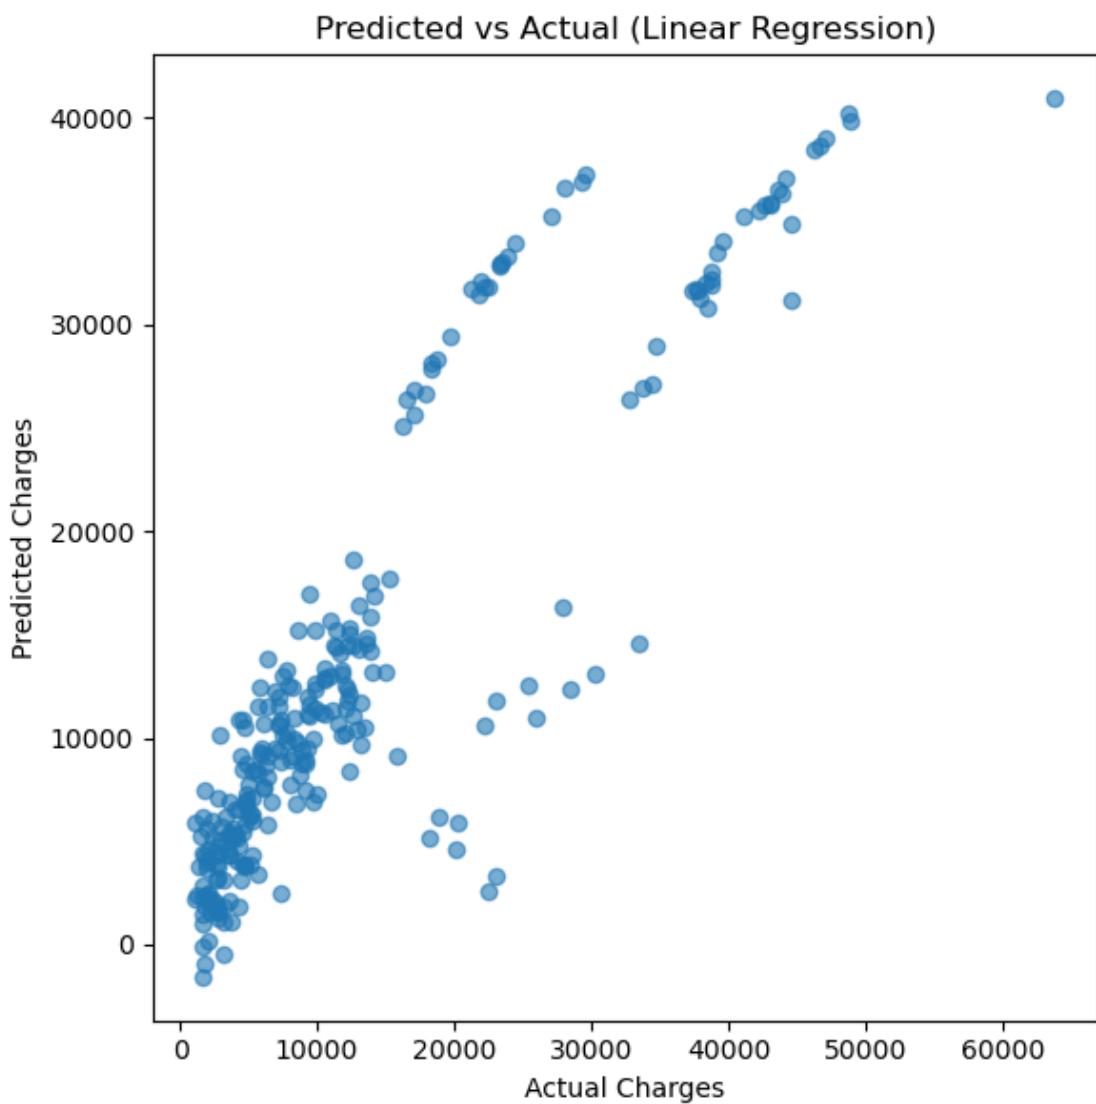
```
[14]: pred_map = {
    "Linear Regression": res_lr["y_pred"],
    "Ridge Regression": res_ridge["y_pred"],
    "Lasso Regression": res_lasso["y_pred"],
    "Ridge (log target)": res_ridge_log["y_pred"],
```

```
}
```

```
y_pred_best = pred_map[best_model]
```

```
[19]: # Plot predicted vs actual to visualize fit quality
```

```
plt.figure(figsize=(6,6))
plt.scatter(y_test, y_pred_best, alpha=0.6)
plt.xlabel("Actual Charges")
plt.ylabel("Predicted Charges")
plt.title(f"Predicted vs Actual ({best_model})")
plt.tight_layout()
plt.savefig(out_dir / "pred_vs_actual.png", dpi=150, bbox_inches="tight")
plt.show()
```



#### 1.5.4 5.4 Residual plot (best model)

Residuals are computed as:

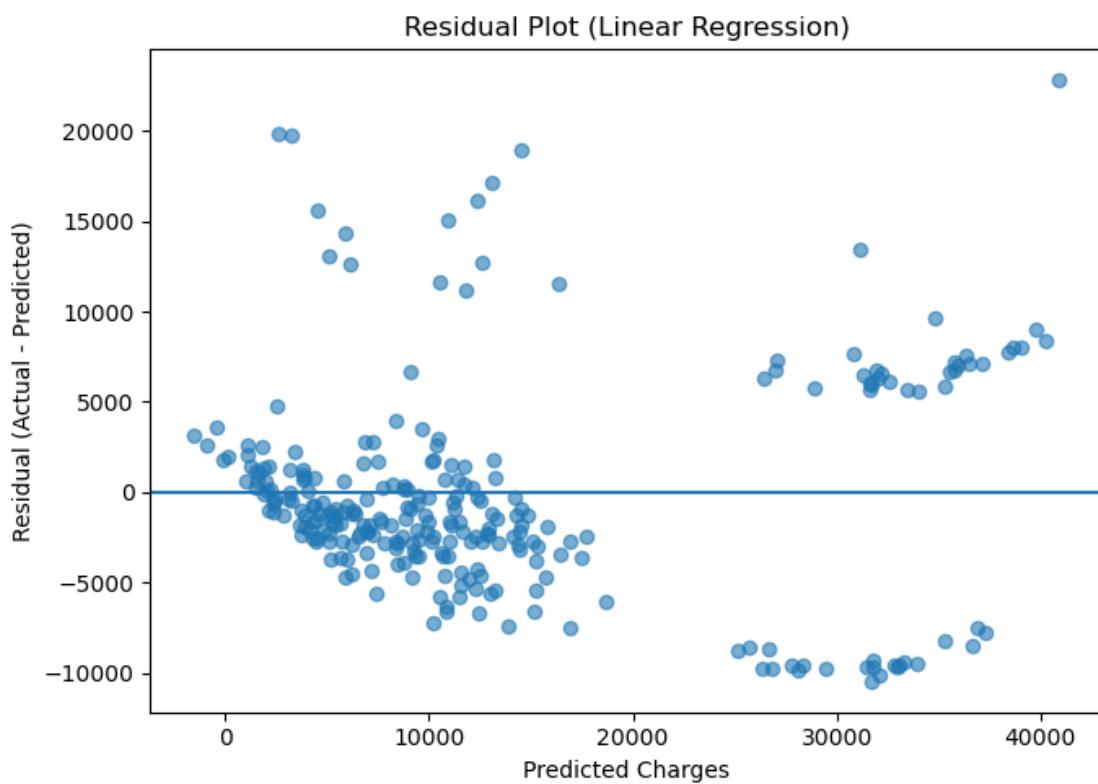
$$\text{Residual} = \text{Actual} - \text{Predicted}$$

A good model should show residuals centered around zero without strong patterns.

```
[20]: # Residual plot to check error patterns and bias

residuals = y_test - y_pred_best

plt.figure(figsize=(7,5))
plt.scatter(y_pred_best, residuals, alpha=0.6)
plt.axhline(0)
plt.xlabel("Predicted Charges")
plt.ylabel("Residual (Actual - Predicted)")
plt.title(f"Residual Plot ({best_model})")
plt.tight_layout()
plt.savefig(out_dir / "residual_plot.png", dpi=150, bbox_inches="tight")
plt.show()
```



### 1.5.5 5.5 Feature importance (Ridge coefficients)

We inspect model coefficients from Ridge Regression to understand which features contribute most to predictions.

Note: Coefficients are influenced by feature scaling and one-hot encoding.

```
[17]: # Feature importance using Ridge coefficients
```

```
model_ridge.fit(X_train, y_train)

ohe = model_ridge.named_steps["preprocess"].named_transformers_["cat"].
    named_steps["onehot"]
cat_names = ohe.get_feature_names_out(categorical_features)

all_feature_names = np.concatenate([numeric_features, cat_names])
coefs = model_ridge.named_steps["model"].coef_

coef_df = pd.DataFrame({
    "feature": all_feature_names,
    "coef": coefs,
    "abs_coef": np.abs(coefs)
}).sort_values("abs_coef", ascending=False)

coef_df.head(15)
```

```
[17]:      feature      coef      abs_coef
5     smoker_no -11791.214744  11791.214744
6     smoker_yes  11791.214744  11791.214744
0        age     3610.436872  3610.436872
1        bmi     2034.082850  2034.082850
2   children     517.098202  517.098202
7  region_northeast  457.758441  457.758441
10 region_southwest -350.175633  350.175633
9   region_southeast -194.321919  194.321919
8  region_northwest   86.739111  86.739111
3     sex_female      7.287392   7.287392
4     sex_male       -7.287392   7.287392
```

### 1.5.6 5.6 Error analysis

To understand failure cases, we inspect the samples with the largest prediction errors. This helps identify where the model struggles most (often extreme high-charge cases).

```
[18]: # Error analysis: largest absolute prediction errors (best model)
```

```
preds_best = pred_map[best_model]

error_df = X_test.copy()
```

```

error_df["y_true"] = np.array(y_test)
error_df["y_pred"] = preds_best
error_df["abs_error"] = np.abs(error_df["y_true"] - error_df["y_pred"])

error_df.sort_values("abs_error", ascending=False).head(10)

```

```
[18]:    age      sex      bmi  children smoker      region      y_true \
543     54  female  47.410         0    yes  southeast  63770.42801
1039    19    male  27.265         2     no  northwest  22493.65964
430     19    male  33.100         0     no  southwest  23082.95533
599     52  female  37.525         2     no  northwest  33471.97189
115     60    male  28.595         0     no  northeast  30259.99556
806     40  female  41.420         1     no  northwest  28476.73499
306     28  female  27.500         2     no  southwest  20177.67113
289     52    male  26.400         3     no  southeast  25992.82104
291     29    male  29.640         1     no  northeast  20277.80751
739     29    male  35.500         2    yes  southwest  44585.45587

              y_pred      abs_error
543  40920.291512  22850.136498
1039  2603.436339  19890.223301
430   3280.691784  19802.263546
599  14560.795906  18911.175984
115  13107.893131  17152.102429
806  12364.784142  16111.950848
306   4574.904104  15602.767026
289  10930.141387  15062.679653
291   5919.186750  14358.620760
739  31161.157390  13424.298480

```

## 1.6 6. Final Analysis

- Linear Regression provides a useful baseline for predicting insurance charges.
- Ridge and Lasso regularization improve stability by controlling coefficient values.
- The log-target Ridge model performed best overall, likely due to reducing the effect of skewed charges.
- Smoking status appears to be the strongest driver of high charges, followed by age and BMI.
- Prediction errors are largest for extreme high-charge cases, suggesting that additional features or non-linear models may improve performance.

[ ]: